

**CSC 7700: Scientific Computing**  
**Module C: Simulations and Application**  
**Frameworks**  
**Lecture 4: Getting Science Out of Computing**  
**Dr. Erik Schnetter**


# Today's Plan

- Review Cactus design principles
- Advanced framework concepts
  - computational efficiency
- Scientific programming
- Simulation data management

# Cactus Design Principles Review

# Basic Framework Design Idea

- Applications consist of many components (*thorns*), glued together by the framework (*flesh*)
- Framework provides main program, components are libraries (control inversion)
- End user assembles code, no central control, easy to add or replace components



# Basic Component Design Idea

- Components are independent entities
- Ideally, a component can stand on its own, without requiring specific other components (higher level and more independent than e.g. C++ class)
- Components need to declare their interfaces (variables, routines, etc.) to the framework

# Some Important Cactus Concepts

- Grid function: container for discretised quantity (e.g. kinetic energy density)
- Thorn list: list of components that should be built into a Cactus configuration (required for building)
- Parameter file: list of components that should be activated, plus their parameter settings (required for running)

# Thorn Specification

Three configuration files per thorn:

- `interface.ccl` declares:
  - an 'implementation' name
  - inheritance relationships between thorns
  - Thorn variables
  - Global functions, both provided and used
- `schedule.ccl` declares:
  - When the flesh should schedule which functions
  - When which variables should be allocated/freed
  - Which variables should be synchronized when
- `param.ccl` declares:
  - Runtime parameters for the thorn
  - Use/extension of parameters of other thorns





# Parallelisation in Cactus

- Implicit parallelisation: No explicit MPI calls in source code, simplifies programming significantly (allows beginners to write parallel code)
- Grid functions provide distributed data structure
- SYNC statements in schedule determine communication (when ghost zones need updating)

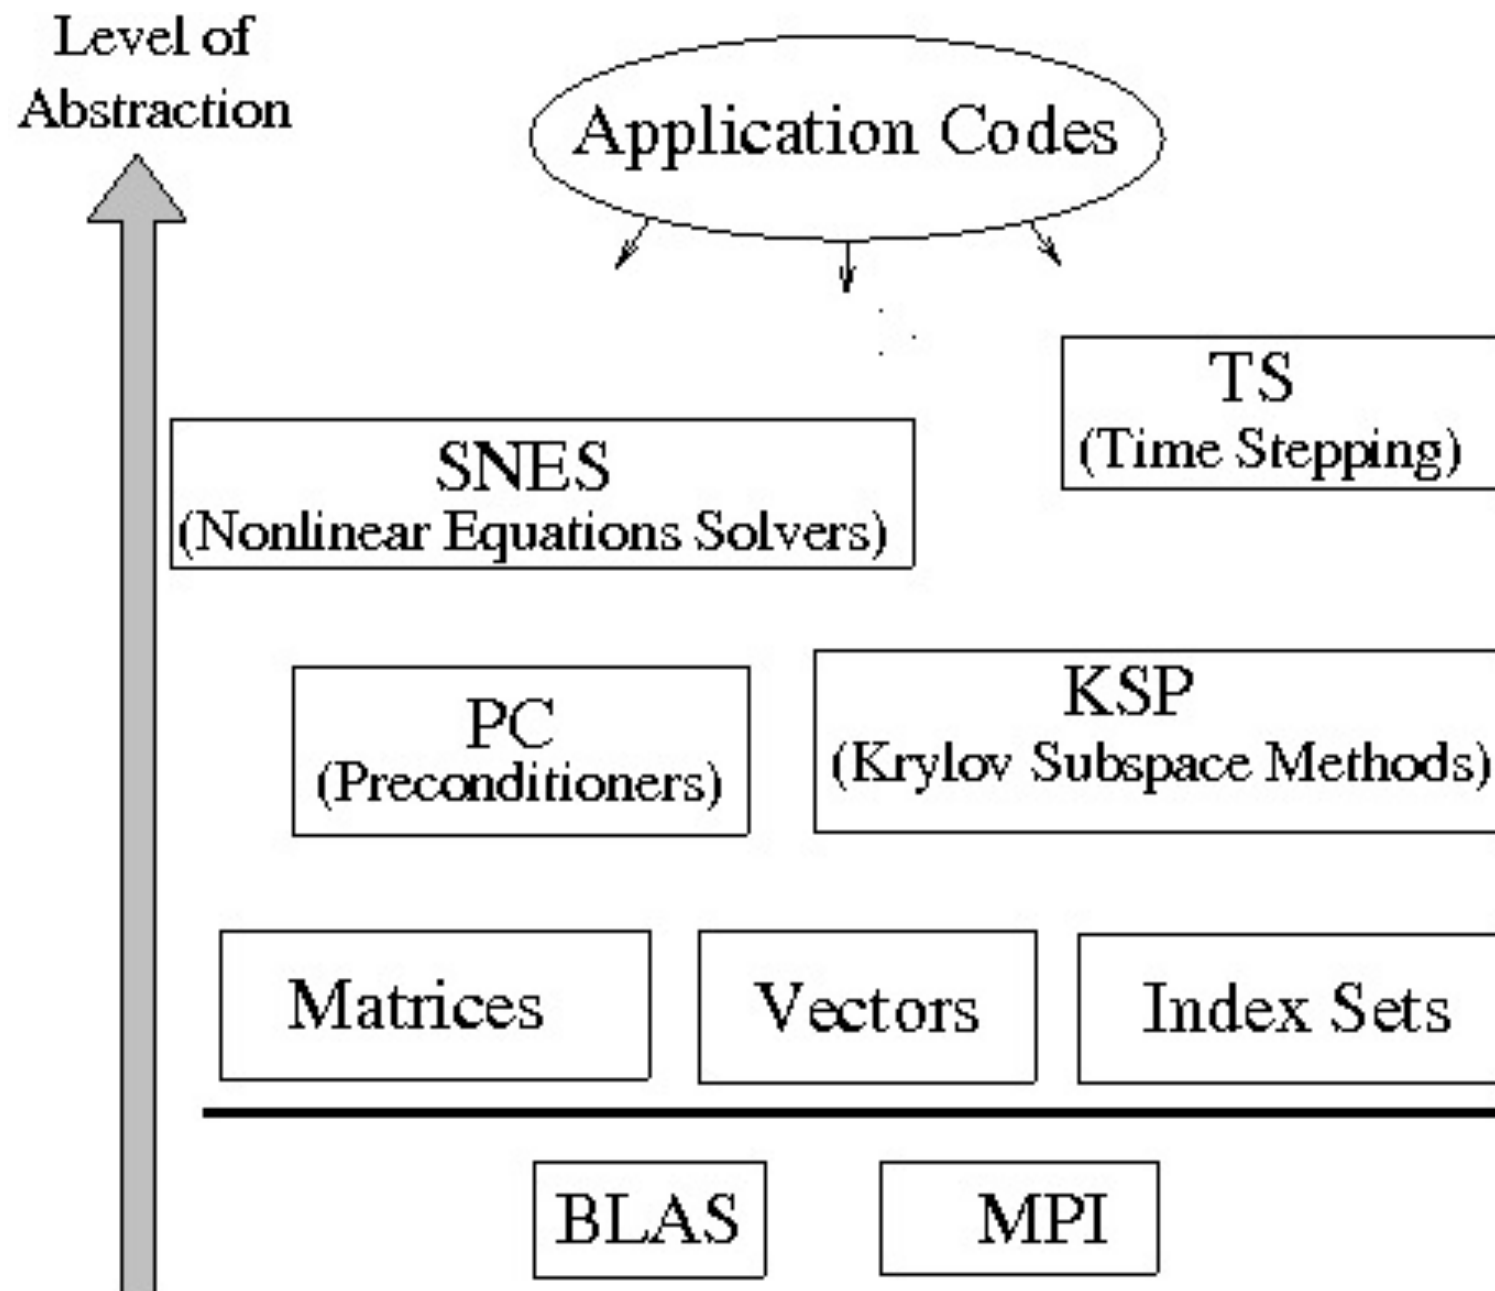
# Questions about Cactus programming?

# Advanced Framework Concepts

# Other Frameworks

- Introducing and comparing several other frameworks for scientific computing
  - PETSc (not a framework)
  - MpCCI (industry standard, commercial)
  - SIERRA (Sandia, not public?)
  - CCA (not alive any more?)

# PETSc (Portable, Extensible Toolkit for Scientific Computation)




- Solver for sparse linear algebra problems
- Also other solvers based on the above
- Infrastructure for sparse linear algebra

**PETSc**

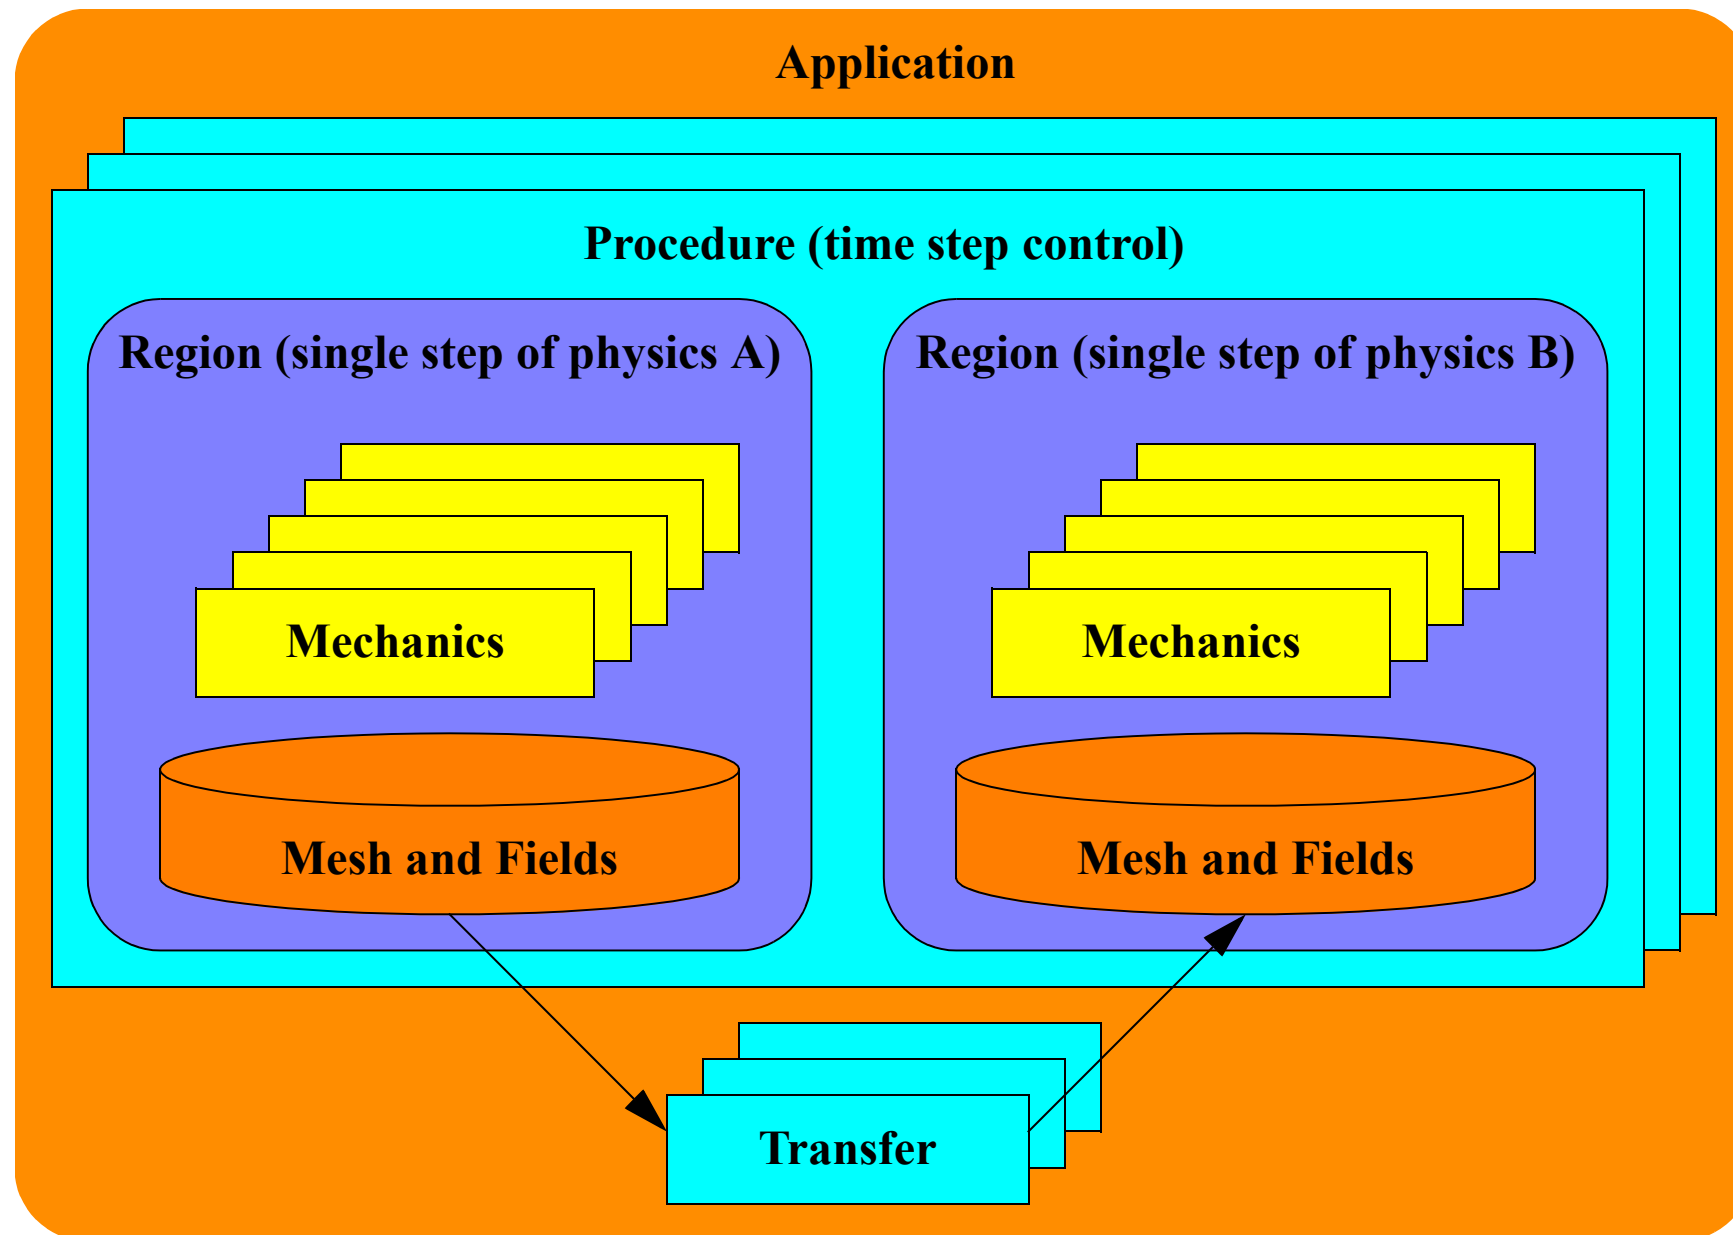
<http://www.mcs.anl.gov/petsc/>

# MpCCI (Multi-physics Code Coupling Interface)



- Can couple arbitrary existing programs (that support some kind of external interface)
- Interface is implemented as additional program

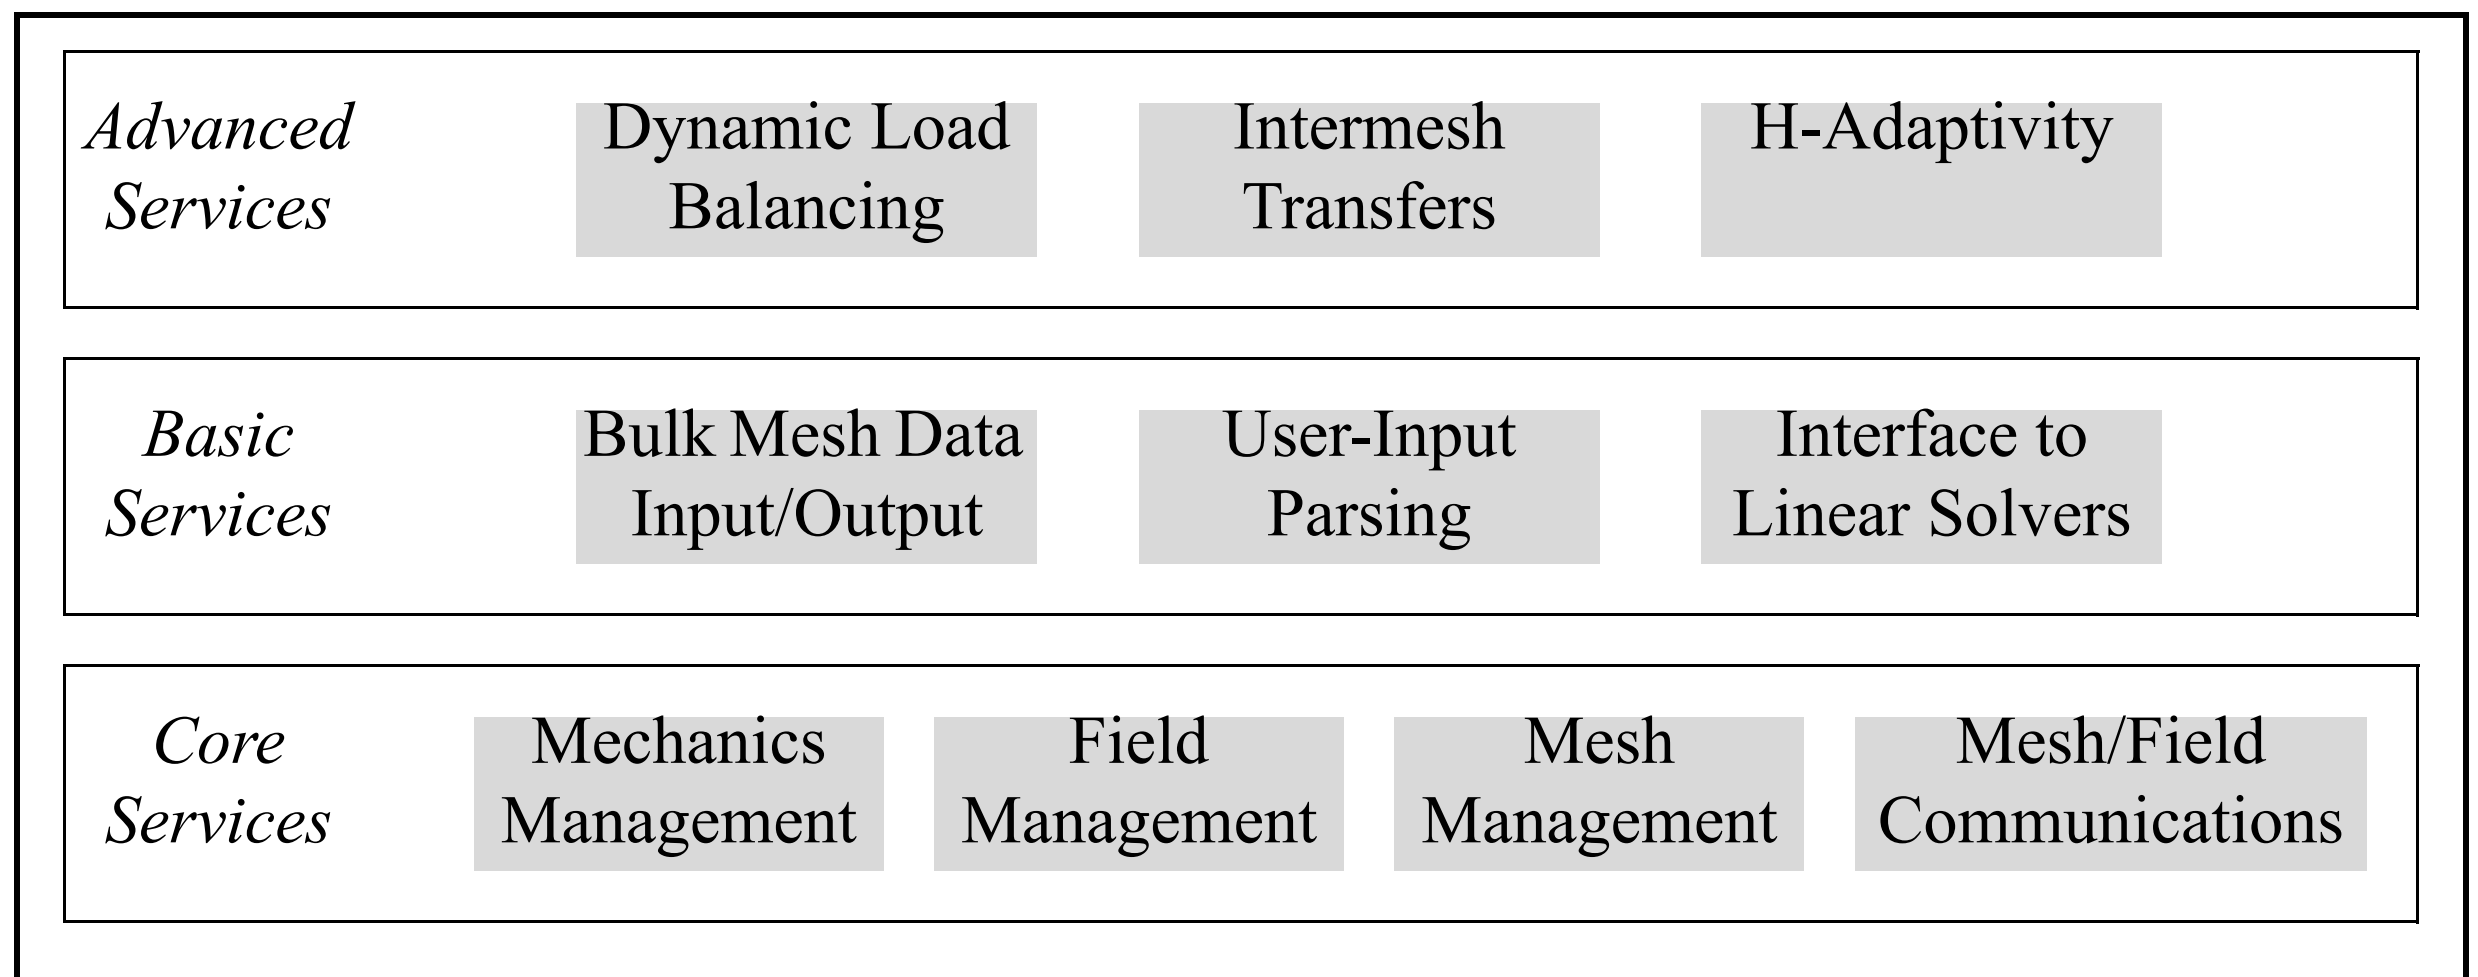
# SIERRA



- Framework targeting Finite Elements (FEM)
- Simulation can have independent regions, coupled via common control mechanism

<http://prod.sandia.gov/techlib/access-control.cgi/2002/023616.pdf>

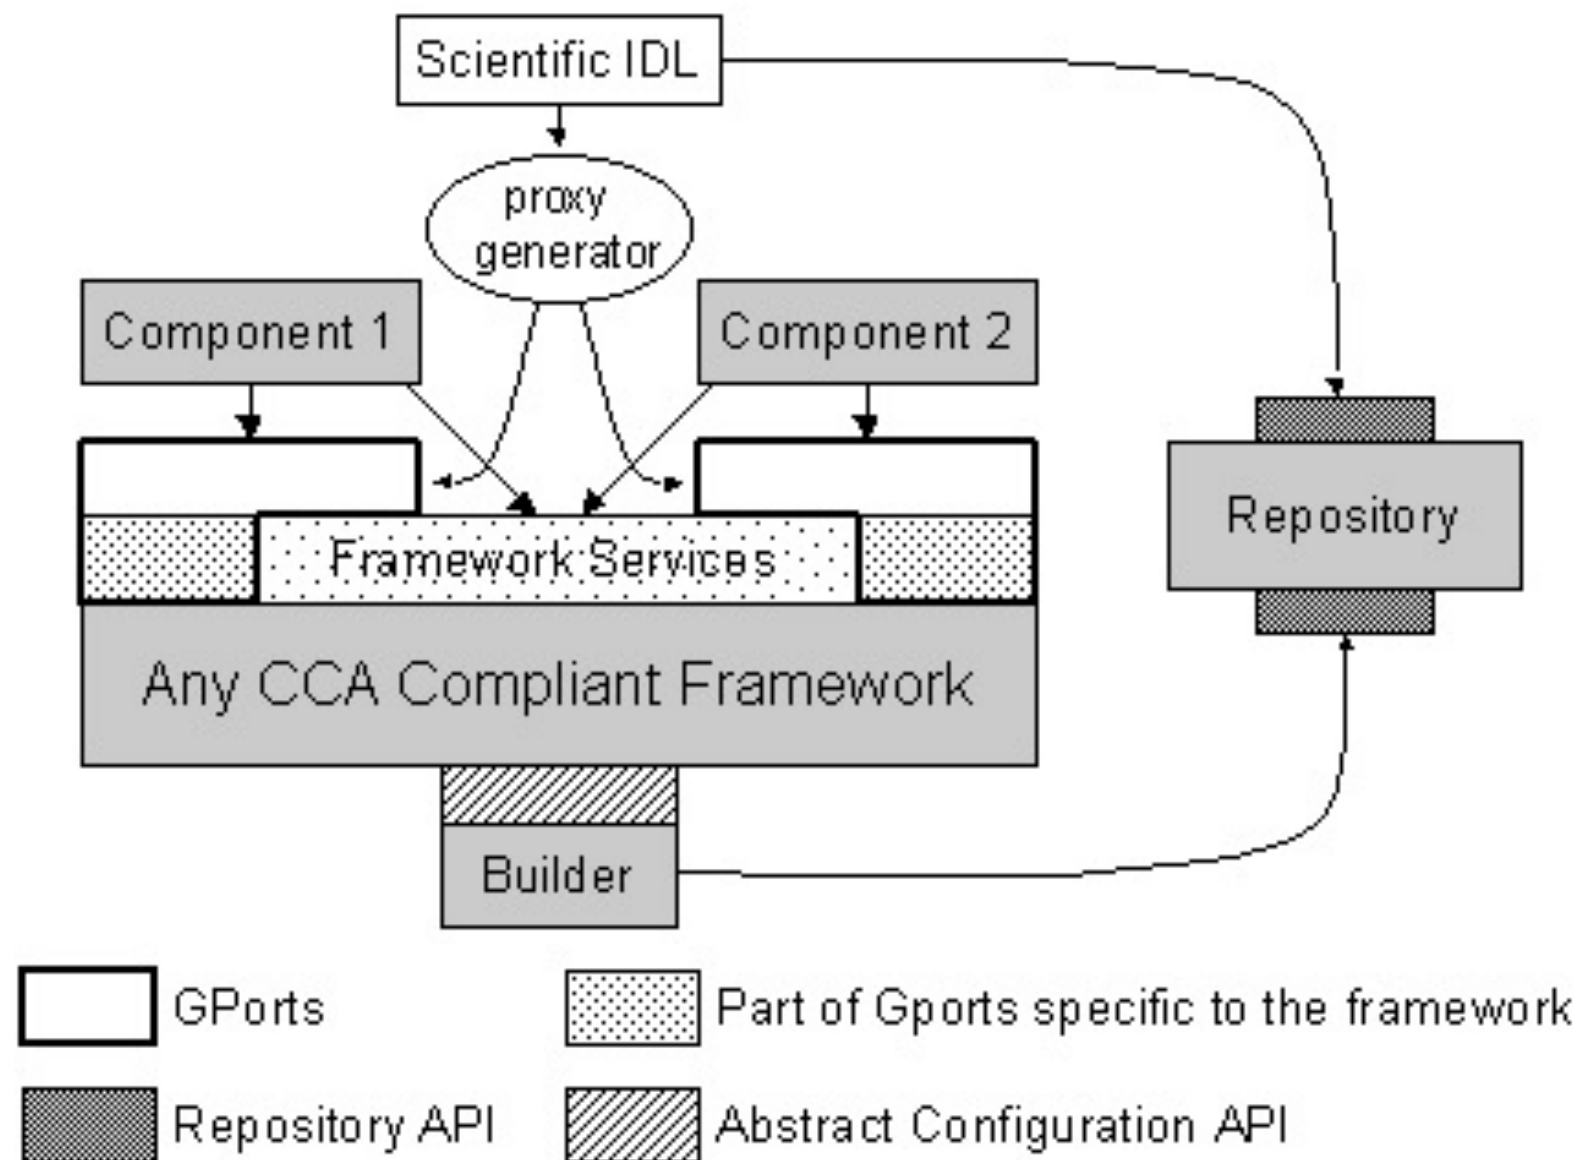
# SIERRA



**Figure 1.2.** SIERRA Framework layered set of services.



# CCA (Common Component Architecture)



- Not a framework, but a specification for framework/component interfaces
- High-level language for defining interfaces: IDL

# Cactus: Driver Thorn

- A *driver* is a special thorn in Cactus that implements parallelism (and memory management)
- Implements “grid function” data type
- This *externalises parallelism*, so that other thorns don’t have to implement parallel algorithms
- However, this places certain restrictions onto other thorns


# Cactus: Driver Thorn

- There must be exactly one driver active (standard Cactus driver is *PUGH*)
- Driver can provide advanced discretisation methods, such as AMR or multi-block (e.g. *Carpet* driver)
- Driver can be based on existing parallel library (e.g. Chombo, Samrai)
- Driver (or related thorn) also provides I/O

# Component Interoperability


[from HPC ASC White Paper]

## Minimal Component Interoperability:




- Physics models are completely uncoupled.
- May exchange static datasets through flat files.

## Shallow Component Interoperability:



- Physics models are loosely coupled.
- Data management and parallelism is independent in each module.
- Exchange common data events via wrappers (web services, etc.).

## Deep Component Interoperability:



- Physics models are tightly coupled.
- Data exchange across shared service infrastructure.

Time

Present

# Application Efficiency

# Data Access

- Simulations handle much data (previous example: 1 billion elements)
- Cannot easily copy data: not enough memory, takes too much time
- Therefore, if possible each process must compute with the data it owns (“bring computation to data”, opposite of a function)

# Data Sharing

- Different components may need to access the same data
  - (Example: your homework project thorn needs to access density, velocity from HydroBase)
- If components are very independent, then data need to be copied
- If data cannot be copied, then the components must interact in some (non-trivial) way

# Component Coupling

- MpCCI: components don't interact, data is exchanged (copied) via "interface process"
- SIERRA (within single region), Cactus: data are stored by framework/driver, different components can access same data
- Components then have to be designed for a specific framework, are not general-purpose any more



# Component Coupling

- No Coupling: independently executing programs, data “sharing” requires copying files
- Loose Coupling: *independent* data management in each component, data sharing requires copying
- Tight Coupling: data are managed outside of components (or by special component), data *sharing is efficient*, but components need to rely on external data manager

# Component Safety

- Efficient data sharing between components requires running in the same address space
- This also means that thorns can (accidentally?) modify each other's data, errors (e.g. array index out of bounds) can propagate between components
- Only compile-time access control and coding standards can provide some safety

# Advanced Framework Concepts Summary

- Many simulation frameworks with many different designs exist
- Fundamental design question is: how tight are components coupled?
- Tight coupling requires shared data management between components
- Trade-off between independence/ease-of-programming/safety and efficiency

# Scientific Programming

# Shared Code Development

- Developing a large code as a group is different from small-scale programming
- There is old code (ten years old) that “belongs to nobody”
- People use “your” code without understanding it
- People make changes to “your” code without understanding it

# Shared Code Development

- Best not to have “your” or “my” code; instead share responsibility
- Everything needs to be programmed defensively, so that wrong usage is always detected
- Functionality requires test cases, so that bad changes can be detected quickly

# Test Cases

- Code can be ten years old and still very good
  - Cannot rewrite old code every year (and introduce new errors every year)
- But need to make sure old code is actually still good, in view of many other changes that may have happened in the mean time
- A *test case* stores program input and expected output, so that any change in behaviour can be detected

# Recovering from Errors

- Mistakes happen, need to be able to undo bad changes
- Important approach: keep complete history of all changes to the code, undo changes when necessary
- Need to use source code management system (e.g. svn); this keeps track of who made when what change



# Working Together

- Source code management system also defines a *single, standard version* of the code, on which everybody is working
- It would be too confusing to send source code around by email, or look into other people's directories
- See tutorials for svn etc.; svn and relatives are indispensable for scientific code development

# Policies

- Working in a group requires policies since programming is a formal task; people need to know what is acceptable
- Coding style (routine names, indentation, commit messages)
- Access rights (using, modifying, adding, committing)
- Testing standards before committing changes
- Peer review before/after making changes


# Component Life Cycle

1. Idea, experimental implementation
2. Prototype, useful for a single paper
3. Production use, more features added, most errors removed, useful for a series of papers
4. Mature, very useful, few changes
5. Outdated, used mostly for historic investigations (still somewhat useful)


# Portability

- Machines become old, outdated, and unreliable after a few years
- HPC systems frequently (once a week?) require maintenance, or are (once a year?) unavailable for a week due to system a upgrade
- Installed software (compilers) may have bugs, takes days or weeks to correct
- Therefore, scientific codes need to be portable, so that one can then quickly use other machines

# New Hardware Architectures



Early multi-core  
revolution graph



# New Hardware Architectures


- Software stays around longer than hardware
  - software: > 15 years (Cactus)
  - hardware: 3 years at most? (Moore's law)
- Software design must not only be portable, but also architecture independent

# GPU Computing

- Increasing the speed of a chip is expensive
- Rule of thumb: Can either have 1 full-speed or 4 half-speed chips for same price
- Multi-core revolution: Every year, put more cores into a computer to make it more powerful
  - Does this also make it faster?

# GPU Computing

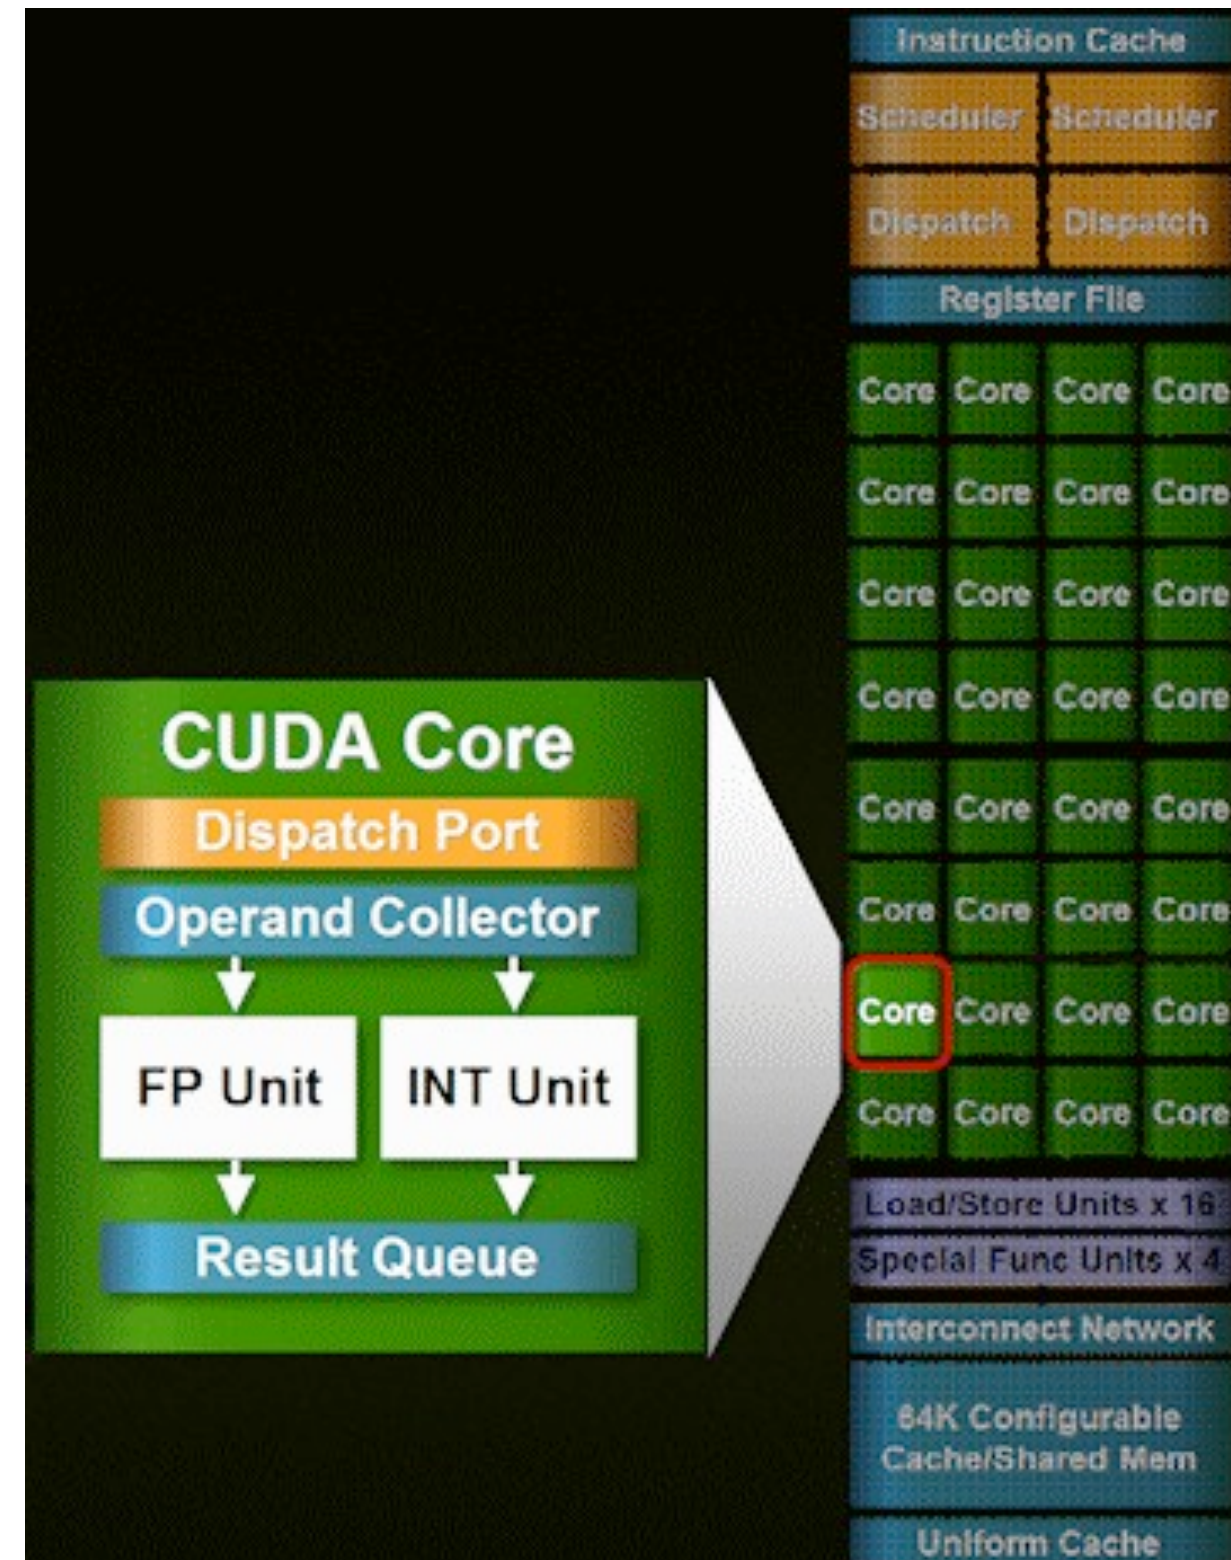
- Started in 2001 with Sony/Toshiba/IBM collaboration for *Cell Processor*
- Has 1 “real” core and 8 slower “additional” cores on a single chip
- Used in Sony’s Playstation





# GPU Computing

- Graphics cards (GPU) have a similar architecture with many cores
- (although these “cores” are, technically, only vector units and not really independent)



# GPU Computing Challenges

- Each core is small and slow, need to use many cores just to be as fast as regular CPU
  - Slower clock speed
  - Less memory per core
- Need to use even more cores to get good speed
- But, if this is possible, then total speed is much larger than that of regular CPUs (of same total cost)

# Framework Architecture Challenges

- Don't want to re-design framework and re-write code for new architectures; framework needs to isolate programmer from architectural changes
- Old architecture (clusters, MPI) has been around for 20 years, but something new is coming now (Roadrunner, first petaflop system in 2008, uses Cell processors)

# Cactus Approach to Architecture Independence

- Separate physics code from computer science code (different thorns, driver)
- Each thorn “sees” only a small part of the overall problem (information hiding)
  - ideally, each physics thorn acts on a single grid point at a time, but that would have too much overhead
  - this externalises parallelism, load balancing, data distribution, data sharing (*largest Cactus success*)
  - this also significantly complicates programming (e.g. calculating total mass) (*largest Cactus problem*)

# Scientific Programming Practices Summary

- Use a source code management system (strongly suggest to do this for homeworks as well)
- Keep track of software versions for each simulation, have test cases to ensure correctness
- Need portability and architecture independence to make programs future-proof

# Simulation Data Management

# Scientific Integrity

- Scientific results must be repeatable
- That is, one must document and preserve the exact software version for each simulation
  - Can either be very careful about this, or use automated tools to help
- Time span is at least months for single paper (review process), years for research groups, difficult for longer times (decades) since hardware changes

# Simulation Data

- Simulations produce much data
- Never enough disk space to keep it all
- Often, auto-clean will delete “old” files
  - see Queen Bee “Work Volume Purge” messages you should be receiving, e.g. this morning



TABLE 1

BINARY SEQUENCES FOR WHICH NUMERICAL SIMULATIONS HAVE BEEN CARRIED OUT, WITH DIFFERENT COLUMNS REFERRING TO THE PUNCTURE INITIAL LOCATION  $\pm x/M$ , THE LINEAR MOMENTA  $\pm p/M$ , THE MASS PARAMETERS  $m_i/M$ , THE DIMENSIONLESS SPINS  $a_i$ , THE NORMALIZED ADM MASS  $\widetilde{M}_{\text{ADM}} \equiv M_{\text{ADM}}/M$  MEASURED AT INFINITY, AND THE NORMALIZED ADM ANGULAR MOMENTUM  $\widetilde{J}_{\text{ADM}} \equiv J_{\text{ADM}}/M^2$ . FINALLY, THE LAST SIX COLUMNS CONTAIN THE NUMERICAL AND FITTED VALUES FOR  $|v_{\text{kick}}|$  (IN km/s),  $a_{\text{fin}}$  AND THE CORRESPONDING ERRORS.


	$\pm x/M$	$\pm p/M$	$m_1/M$	$m_2/M$	$a_1$	$a_2$	$\widetilde{M}_{\text{ADM}}$	$\widetilde{J}_{\text{ADM}}$	$ v_{\text{kick}} $	$ v_{\text{kick}}^{\text{fit}} $	err. (%)	$a_{\text{fin}}$	$a_{\text{fin}}^{\text{fit}}$	err. (%)
$r0$	3.0205	0.1366	0.4011	0.4009	-0.584	0.584	0.9856	0.825	261.75	258.09	1.40	0.6891	0.6883	0.12
$r1$	3.1264	0.1319	0.4380	0.4016	-0.438	0.584	0.9855	0.861	221.38	219.04	1.06	0.7109	0.7105	0.06
$r2$	3.2198	0.1281	0.4615	0.4022	-0.292	0.584	0.9856	0.898	186.18	181.93	2.28	0.7314	0.7322	0.11
$r3$	3.3190	0.1243	0.4749	0.4028	-0.146	0.584	0.9857	0.935	144.02	146.75	1.90	0.7516	0.7536	0.27
$r4$	3.4100	0.1210	0.4796	0.4034	0.000	0.584	0.9859	0.971	106.11	113.52	6.98	0.7740	0.7747	0.08
$r5$	3.5063	0.1176	0.4761	0.4040	0.146	0.584	0.9862	1.007	81.42	82.23	1.00	0.7948	0.7953	0.06
$r6$	3.5988	0.1146	0.4638	0.4044	0.292	0.584	0.9864	1.044	45.90	52.88	15.21	0.8150	0.8156	0.07
$r7$	3.6841	0.1120	0.4412	0.4048	0.438	0.584	0.9867	1.081	20.59	25.47	23.70	0.8364	0.8355	0.11
$r8$	3.7705	0.1094	0.4052	0.4052	0.584	0.584	0.9872	1.117	0.00	0.00	0.00	0.8550	0.855	0.00
$ra0$	2.9654	0.1391	0.4585	0.4584	-0.300	0.300	0.9845	0.8250	131.34	132.58	0.95	0.6894	0.6883	0.16
$ra1$	3.0046	0.1373	0.4645	0.4587	-0.250	0.300	0.9846	0.8376	118.10	120.28	1.85	0.6971	0.6959	0.17
$ra2$	3.0438	0.1355	0.4692	0.4591	-0.200	0.300	0.9847	0.8499	106.33	108.21	1.77	0.7047	0.7035	0.17
$ra3$	3.0816	0.1339	0.4730	0.4594	-0.150	0.300	0.9848	0.8628	94.98	96.36	1.46	0.7120	0.7111	0.13
$ra4$	3.1215	0.1321	0.4757	0.4597	-0.100	0.300	0.9849	0.8747	84.74	84.75	0.01	0.7192	0.7185	0.09
$ra6$	3.1988	0.1290	0.4782	0.4602	0.000	0.300	0.9850	0.9003	63.43	62.19	1.95	0.7331	0.7334	0.04
$ra8$	3.2705	0.1261	0.4768	0.4608	0.100	0.300	0.9852	0.9248	41.29	40.55	1.79	0.7471	0.7481	0.13
$ra10$	3.3434	0.1234	0.4714	0.4612	0.200	0.300	0.9853	0.9502	19.11	19.82	3.72	0.7618	0.7626	0.11
$ra12$	3.4120	0.1209	0.4617	0.4617	0.300	0.300	0.9855	0.9750	0.00	0.00	0.00	0.7772	0.7769	0.03
$s0$	2.9447	0.1401	0.4761	0.4761	0.000	0.000	0.9844	0.8251	0.00	0.00	0.00	0.6892	0.6883	0.13
$s1$	3.1106	0.1326	0.4756	0.4756	0.100	0.100	0.9848	0.8749	0.00	0.00	0.00	0.7192	0.7185	0.09
$s2$	3.2718	0.1261	0.4709	0.4709	0.200	0.200	0.9851	0.9251	0.00	0.00	0.00	0.7471	0.7481	0.13
$s3$	3.4098	0.1210	0.4617	0.4617	0.300	0.300	0.9855	0.9751	0.00	0.00	0.00	0.7772	0.7769	0.03
$s4$	3.5521	0.1161	0.4476	0.4476	0.400	0.400	0.9859	1.0250	0.00	0.00	0.00	0.8077	0.8051	0.33
$s5$	3.6721	0.1123	0.4276	0.4276	0.500	0.500	0.9865	1.0748	0.00	0.00	0.00	0.8340	0.8325	0.18
$s6$	3.7896	0.1088	0.4002	0.4002	0.600	0.600	0.9874	1.1246	0.00	0.00	0.00	0.8583	0.8592	0.11
$t0$	4.1910	0.1074	0.4066	0.4064	-0.584	0.584	0.9889	0.9002	259.49	258.09	0.54	0.6868	0.6883	0.22
$t1$	4.0812	0.1103	0.4062	0.4426	-0.584	0.438	0.9884	0.8638	238.37	232.62	2.41	0.6640	0.6658	0.27
$t2$	3.9767	0.1131	0.4057	0.4652	-0.584	0.292	0.9881	0.8265	200.25	205.21	2.48	0.6400	0.6429	0.45
$t3$	3.8632	0.1165	0.4053	0.4775	-0.584	0.146	0.9879	0.7906	174.58	175.86	0.73	0.6180	0.6196	0.26
$t4$	3.7387	0.1204	0.4047	0.4810	-0.584	0.000	0.9878	0.7543	142.62	144.57	1.37	0.5965	0.5959	0.09
$t5$	3.6102	0.1246	0.4041	0.4761	-0.584	-0.146	0.9876	0.7172	106.36	111.34	4.68	0.5738	0.5719	0.33
$t6$	3.4765	0.1294	0.4033	0.4625	-0.584	-0.292	0.9874	0.6807	71.35	76.17	6.75	0.5493	0.5475	0.32
$t7$	3.3391	0.1348	0.4025	0.4387	-0.584	-0.438	0.9873	0.6447	35.36	39.05	10.45	0.5233	0.5227	0.11
$t8$	3.1712	0.1419	0.4015	0.4015	-0.584	-0.584	0.9875	0.6080	0.00	0.00	0.00	0.4955	0.4976	0.42
$u1$	2.9500	0.1398	0.4683	0.4685	-0.200	0.200	0.9845	0.8248	87.34	88.39	1.20	0.6893	0.6883	0.15
$u2$	2.9800	0.1384	0.4436	0.4438	-0.400	0.400	0.9846	0.8249	175.39	176.78	0.79	0.6895	0.6883	0.17
$u3$	3.0500	0.1355	0.3951	0.3953	-0.600	0.600	0.9847	0.8266	266.39	265.16	0.46	0.6884	0.6883	0.01
$u4$	3.1500	0.1310	0.2968	0.2970	-0.800	0.800	0.9850	0.8253	356.87	353.55	0.93	0.6884	0.6883	0.01

# Simulation Laboratory

- Publications usually based on multiple simulations
  - ...performed by different people at different times on different systems...
  - ...with different versions of the code...
- Tedious to keep track of everything, administrative burden

# Data Curation

- Simulations are expensive (\$1M for previous table?)
- Scientific results should be “eternal”, must (?) be able to go back and re-analyse data even after years
  - Can you?
- What if project is interesting and takes years to complete?
  - How long will your class projects survive?
  - If they don't, then why are they useful?




# Finding Data

- If you indeed keep all old school programming projects around, how do you find something later if you need it?
- How do you find simulation data?
  - If machines are shut down after several years?
  - If they were performed by others who left the group years ago?
  - Do you just spend another \$1M to re-run the simulations? Does the old code still work? (How much does it cost to archive the data? To keep the code working?)

# Summary

# Simulation Science Basics

- Physics model described by PDEs
- PDEs discretised in space and time, leading to many (billions) of elements
- Discrete equations implemented on large supercomputers
- A simulation iterates many (millions of) times to produce its result



# Simulating Complex Systems

- Require parallelism for efficient executing
  - These days, most supercomputers have cluster architecture programmed with MPI
  - Domain decomposition is the prevalent strategy for parallelism
- Real-world simulation codes are large, framework model can provide good architecture for this

# Framework Architecture

- A framework provides glue to combine independent components
- This allows large-scale, loosely-coupled program development, ideal for today's international research groups
- Cactus is such a framework
  - Components need to declare their interfaces to the framework



# Getting Science Out of Computing

- Components can be loosely or tightly coupled, depending on independence/efficiency trade-off
- On HPC systems, data placement defines efficiency since moving data is expensive
- Consider new/future system architectures
- Scientific integrity requires archiving simulation data