

**CSC 7700: Scientific Computing**  
**Module C: Simulations and Application**  
**Frameworks**  
**Lecture I: Simulation Science Basics**  
**Dr. Erik Schnetter**

# Goal







- This module *Simulations and Application Frameworks* will teach
  - how a typical simulation code looks like,
  - how it is used in practice (by physicists or engineers),
  - and what some of the major concerns in such a code are.

# Literature

- Article: T. Goodale et al., *The Cactus Framework and Toolkit: Design and Applications*
- Book: M.T. Heath, *Scientific Computing: An Introductory Survey*
- Documentation: Cactus Users' Guide
- Tutorial: Einstein Toolkit Tutorial
  - Details see wiki page

# From Physics to Simulation


# Simulations



# Why Use Simulations?

- Flame propagation in combustion engine: *understand* behaviour that is too fast or too small
- Hurricane modelling: *predict* behaviour
- Car crash testing: *engineer* better devices
- Video games: *create* a fantasy world similar to the real one

Laws of Physics  
(or Chemistry, Biology, ...)



Simulation

**Laws of Physics  
(or Chemistry, Biology, ...)**



**Mathematics**




**Supercomputers**



**Simulation**





- The physics that is to be simulated is expressed in “the language of Mathematics”
- Called *Scientific Computing* or *Numerical Analysis*

- The resulting systems of equations are solved on large computers
- Called *Supercomputers* because they are as large and awkward as a supertanker




# Systems and Equations

- The state of a system is described via variables (density, velocity, pressure, etc.)
- Laws of Physics can then often be described via *PDEs* (Partial Differential Equations)
- A PDE describes how a system is *changing* depending on its current *state*



# PDE Example

Euler equation:  
(hydrodynamics) 
$$\frac{Dv}{Dt} = -\frac{1}{\rho} \frac{\partial p}{\partial x}$$




- $\rho$ : density,  $p$ : pressure,  $v$ : velocity;  
 $t$ : time,  $x$ : position
- Interpretation:  
Consider a small chunk of matter. If the pressure to its left and right are different ( $\partial p$ ), then it will be accelerated ( $Dv$ ).
- This assumes a co-moving coordinate system (Lagrangian picture). In a fixed coordinate system, additional terms will appear in the Euler equation.

# PDE Interpretation

- A *PDE* (partial differential equation) tells us how a system is changing, if we know the state of the system
- Starting from an *Initial Condition*, we can thus simulate the behaviour of a system by meticulously tracking how the system is changing

# Discretisation



- PDEs describe continuum systems (car body, water, air); these have infinitely many degrees of freedom
- Reduce complexity by approximation via a *discrete system* instead
- Compare e.g. pixels on a TV screen, surface triangulation for visualisation
- Many possibilities:
- finite elements (e.g. small rigid triangles)
- finite volumes (e.g. small cubes)
- finite differences (sample solution on regular grid)
- particles (small chunks of matter)
- many more...


# Discretisation Error

- Discretising is an approximation and thus leads to an error
- Can use a finer discretisation (higher *resolution*) to reduce this error
- *Order of Accuracy* describes how this error scales with the resolution , e.g.  
fourth order:  $E = O(h^4)$   
doubling resolution reduces error by 16




# Simulation Procedure

- Choose PDE that describes system well
- Discretise PDE
- Set up initial condition
- Follow each element of the system over many many tiny steps
- A simulation can have billions of elements with millions of steps, taking weeks of computing time



# Caveat

- Some systems are described not by PDEs but otherwise (e.g. coupled ODEs, discrete transitions)
- Sometimes not time evolution is interesting, but e.g. equilibrium configuration
- Usually (in real life), PDEs and initial conditions are only *approximations or guesses*, and simulation results *may not be reliable*





# Connection to Other Modules

- Some systems are not described by PDEs:  
**Distributed Scientific Computing**
- Simulations produce large output files  
(billions of elements, millions of steps):  
**Networks and Data**
- To understand results, need to “undo”  
formulation as PDE and discretisation:  
**Scientific Visualisation**


# Ingredients of a Simulation

# Key Concepts in a Numerical Simulation

- Simulation Domain: the part of the world that is simulated, often just a small box
- Resolution: the accuracy of the discretisation; higher is better (and more expensive)
- Evolution System: (discretised) PDE system
- Initial Condition: initial state
- Boundary Condition: what to do at the (artificial?) domain boundaries, often also PDEs
- Output Variables: which part of the solution should be output -- it is often too expensive to output everything

# WaveToy Thorn: Wave Equation

For a given source function  $S(x, y, z, t)$  find a scalar wave field  $\varphi(x, y, z, t)$  inside the domain  $\mathcal{D}$  with a boundary condition:



- inside  $\mathcal{D}$ :

$$\frac{\partial^2 \varphi}{\partial t^2} = c^2 \Delta \varphi + S$$

- on the boundary  $\partial\mathcal{D}$ :


$$\varphi|_{\partial\mathcal{D}} = \varphi(t = 0)$$



# WaveToy Thorn: Discretization

Discretization:

approximating continuous function  $\varphi(x, t)$  with a grid function  $\varphi_i^{(n)}$ :



$$\frac{\partial^2 \varphi}{\partial t^2} = c^2 (\partial_x^2 \varphi) + S$$

$$\Downarrow (c \equiv 1)$$

$$\frac{\varphi_i^{(n+1)} - 2\varphi_i^{(n)} + \varphi_i^{(n-1)}}{2\Delta t^2} = \frac{\varphi_{i+1}^{(n)} - 2\varphi_i^{(n)} + \varphi_{i-1}^{(n)}}{2\Delta x^2} + S_i^{(n)}$$




# Basic Structure of a Simulation Code

- State (solution) stored in large “vectors”
- Routine to set up initial condition
- Routine to perform many identical steps (applying discretised PDE or similar)
- I/O methods to write solution to disk
- Run as batch job without user interaction

# Storing the Solution

- After discretising a PDE, one obtains many (...billions...) very similar elements (cells, points, particles, ...)
- Best handled in efficient container structure: Fortran array, C++ vector, etc. (maybe also tree structure)
- Code contains many constructs that iterate over these elements



# Initial Condition

- Often generated by external method, then read in from file
- Can also checkpoint, and then restart where previous simulation left off
- Initial data can be large; example:
  - 1 billion elements,
  - 5 variables/element,
  - 8 Byte/variable: total 40 GByte



# Parallel Computing

- Cannot store solution on a single node; parallel programming via MPI is a must
- These days, only Fortran, C and C++ are viable languages for programming a supercomputer
- There is research in other, simpler ways, e.g. Unified Parallel C, Co-Array Fortran, or ParalleX (here at LSU)

# (Time) Stepping

- Performing many identical steps to arrive at the solution
- Simulations can take long; example:
  - 1 billion elements,
  - 1000 Flop/element per step,
  - 1 million steps,
  - CPU speed 10 GFlop/sec:
  - total 28,000 CPU hours (3.2 CPU years),
  - or 12 days when running on 100 CPUs

# Batch Processing

- Since simulations take so long, cannot supervise them manually
  - Cannot be awake at all times
  - Each user error can destroy weeks of data
  - Supercomputers are expensive; cannot wait for the next user input
- Need to *plan* simulations carefully ahead of time, then let them run automatically

# Batch Processing 2

- Need to *plan* simulations carefully ahead of time, then let them run automatically...
- ...so that each error is only discovered weeks later!
- Using a supercomputer thus requires much expertise and experience, patience, and a high tolerance for frustration
- **This points to a large problem in supercomputer usability these days**

# Ingredients of a Simulation

- Many simulation programs have a similar structure
- This structure is determined by the physics description (PDEs and discretisation)
- A simulation handles many small elements of data, and iterates over them many times

# Ingredients of a Supercomputer

# Fast vs. Large


- Supercomputers are not fast, they are large
- They are not interactive (like a notebook or workstation), they operate in batch mode
- Their hardware is complex -- I am going to describe the user's point of view only here



# Remote Access



- Supercomputers are located in far away places, need to use ssh/gsissh to access
- Log in is to *front end (head node)* only, usually a large workstation
- Cannot (or should not) use front end to run simulations





# File Systems




- Supercomputers need large file systems to store simulation data, often many 100 TByte
- For management and performance reasons, usually split into different parts with different properties
- Different on each supercomputer -- read documentation!
- Home directory: GBytes per user, many small files, backed up
- Data directory: TBytes per user, few large files, backed up, tape backend
- Scratch directory: no quota, few large files, often automatically deleted

# Compute Nodes, Interconnect



- Most supercomputers have a *cluster* architecture with many compute nodes
- Each node has (4 to 32?) cores, similar to a large workstation
- Nodes are connected via a low-latency *communication network* (e.g. Infiniband)
- Overall system has (128 to >8,000?) nodes, or up to 100k cores
- My personal scale:
  - <1k cores: small,
  - <10k cores: medium
  - >10k cores: large

# Batch System



- Cannot (or should not) use compute nodes directly
- Need to submit *job* to *batch system*, requesting  $N$  nodes...
- ... wait (a few days?) ...
- ... then the job runs
- (... and then one discovers one's errors)
- There is a run time limit, often 24h or 48h
- ... which is inconvenient if one needs to run for 2 weeks: checkpoint/restart
- Batch systems ensure that a supercomputer is not idle; there are always jobs waiting to be executed



# Allocations



- Need to ensure fair use of supercomputer, prevent individual users from monopolising it
- Typically, an *allocation process* decides who can use how much of a supercomputer's time during a year (similar to writing a grant proposal)
- 1 CPU hour costs about 5 cents (10 cents on Amazon ECC)
- With this metric, Queen Bee produces about \$270 worth of CPU time every hour

# Software

- Installed/available software is system dependent, not just standard Unix systems
- Therefore cannot just install binaries, need to build software manually (or ask administrators to do that)
- HPC developers often prefer command line tools, don't use GUIs (which may not be available)
- (But: Eclipse and PTP may change this)



# Ingredients of a Supercomputer

- Obtaining an *Allocation*,
- Logging in to a *Front End*,
- Submitting jobs to a *Batch System*,
- Simulation executes on *Compute Nodes* connected via a *Communication Network*,
- Storing data in various *File Systems*.

# Sample Session: Einstein Toolkit

# Tutorial

- The tutorial instructions are at [http://docs.einsteintoolkit.org/et-docs/Tutorial\\_for\\_New\\_Users](http://docs.einsteintoolkit.org/et-docs/Tutorial_for_New_Users)
- Note: These instructions require an *account* and an *allocation* on Queen Bee. Obtaining these may take several business days!



# Homework

- Follow these instructions. Skip the “Additional ...” parts.
- Write a report detailing how many cores the simulation used, how much CPU time it required, and how much disk space the simulation output occupies.
- State which allocation you used, and how long you needed to wait in the queue.
- Include the gnuplot graphs in your report.